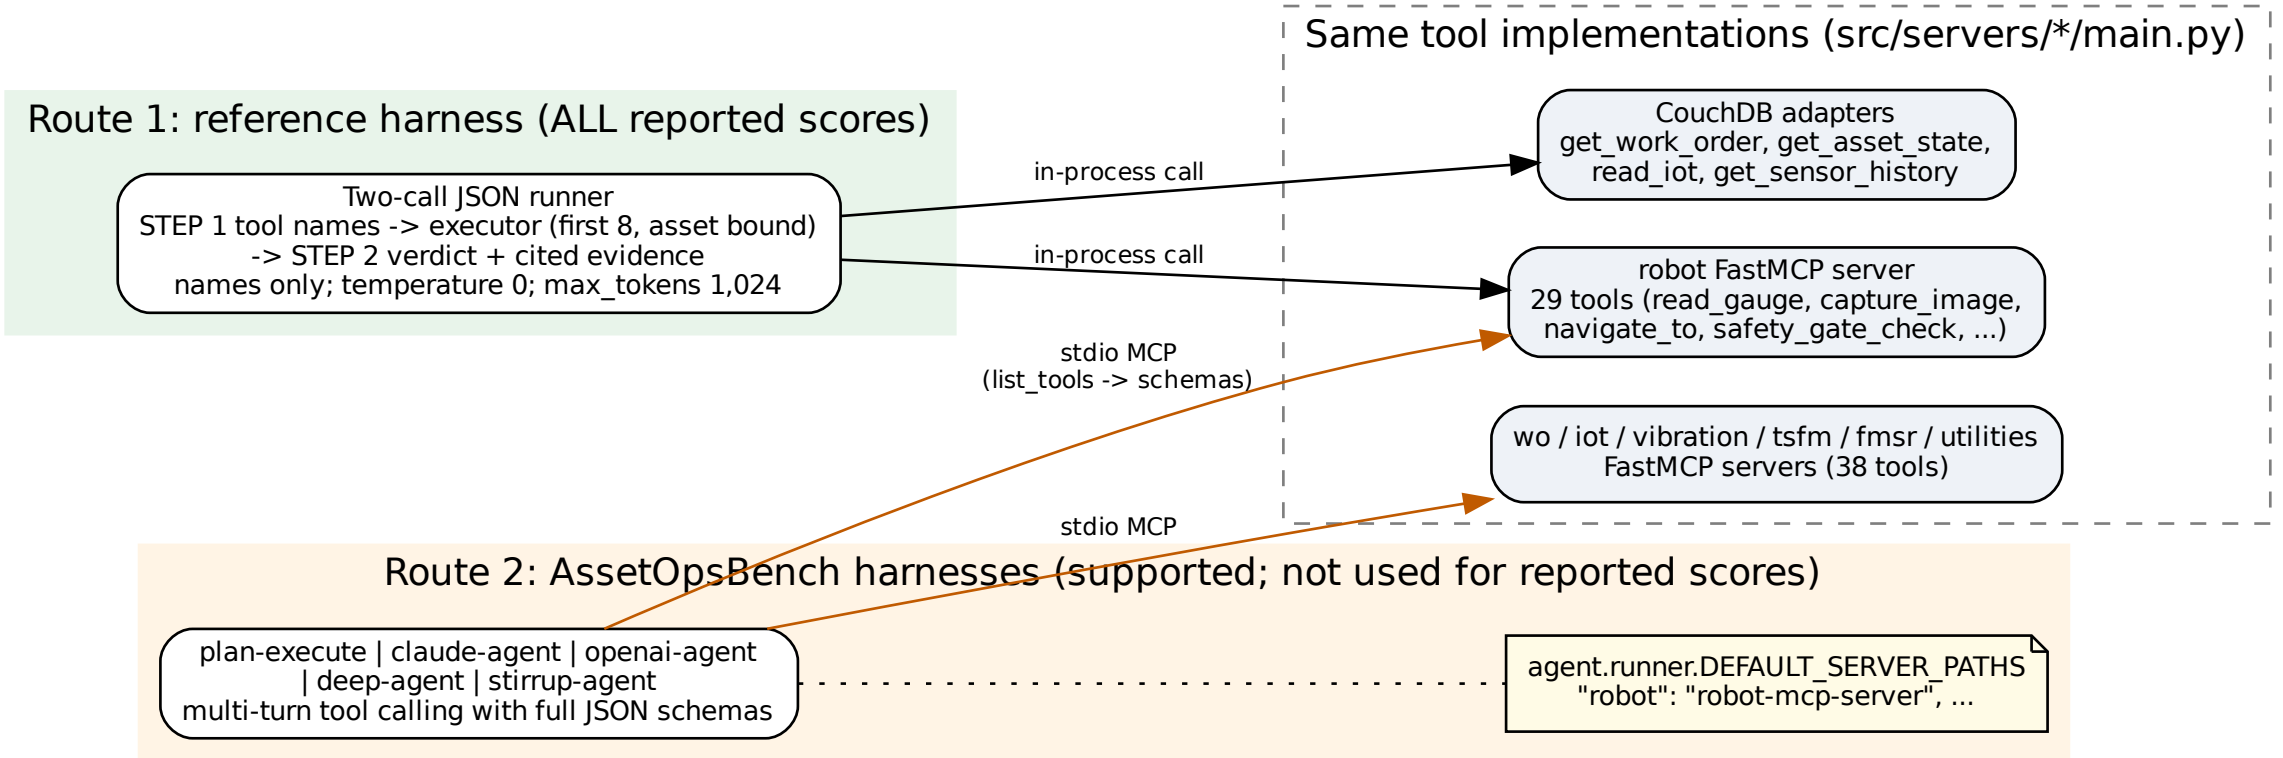
Route 1: reference harness (ALL reported scores)
Two-call JSON runner
STEP 1 tool names -> executor (first 8, asset bound)
-> STEP 2 verdict + cited evidence
names only; temperature 0; max_tokens 1,024
in-process call
in-process call
Same tool implementations (src/servers/*/main.py)
CouchDB adapters
get_work_order, get_asset_state,
read_iot, get_sensor_history
robot FastMCP server
29 tools (read_gauge, capture_image,
navigate_to, safety_gate_check, ...)
wo / iot / vibration / tsfm / fmsr / utilities
FastMCP servers (38 tools)
stdio MCP
(list_tools -> schemas)
stdio MCP
Route 2: AssetOpsBench harnesses (supported; not used for reported scores)
plan-execute | claude-agent | openai-agent
| deep-agent | stirrup-agent
multi-turn tool calling with full JSON schemas
agent.runner.DEFAULT_SERVER_PATHS
"robot": "robot-mcp-server", ...